

ShieldCall: Streaming Detection of Vishing Scripts and Vcoded Speech under Telephone Channels

Vubangsi Mercel

Independent researcher vmercel@gmail.com

Abstract

Telephone fraud combines social-engineering language with increasingly accessible synthetic or vocoded speech. Most published countermeasures treat the two problems separately, and many acoustic systems are evaluated on clean or lightly augmented audio rather than on the narrowband, quantized, packet-loss conditions of live calls. This paper describes ShieldCall, a streaming software stack that scores linguistic fraud-intent and acoustic production artifacts on a shared 8 kHz timeline and fuses them with explicit disagreement rules. We do not claim state-of-the-art equal error rate on ASVspoof. On a held-out set of author-written call scripts, a stage tracker that models vishing as a progression of discourse stages reaches AUC 0.88, while a high-precision keyword baseline falls to AUC 0.42. On speaker-disjoint Mini LibriSpeech, a pulse-formant vocoder is separated from bona fide speech at 8 kHz after bandlimiting, whereas a low-order LPC vocoder is not: residual features are at chance under narrowband (AUC 0.49). When a call is labeled as a threat if it is either a scam script or vocoded, rule-based fusion raises disagreement-cell recall at a 0.5 threshold from 0.30 (naive sum) to 1.00, at the cost of a higher false-positive rate on safe calls. We further pose a timing question that utterance-level scores cannot ask: is a production *change-point* closer to a harvest-class stage than a circular-shift null would allow? The statistic ranks perfectly on synthetic point processes (AUC 1.00) and does *not* on pulse-formant splices into Mini LibriSpeech (AUC 0.47, $n=6$ pairs). Adaptive conformal inference of Gibbs and Candès tracks a 0.90 target with empirical coverage 0.885 on a labeled score stream with a mid-course prevalence shift. The code is released. Do not cite the audio handoff number as a detection success.

Keywords: vishing, voice spoofing, telephone channel, streaming detection, score fusion.

1 Introduction

Impersonation and payment-harvesting calls remain a large consumer-protection problem in the United States [1]. Two technical facts make the problem harder than keyword spotting on a clean recording. First, the incriminating evidence is often the *script*: a human speaker can walk a victim through authority, urgency, credential harvest, and prepaid-card payment without ever using a neural vocoder. Second, when a voice is synthetic, the

cues that laboratory anti-spoofing systems rely on are distorted by μ -law quantization, 300–3400 Hz bandlimiting, and packet-loss concealment [2, 3].

ShieldCall is a software system for the joint, streaming case. Audio is framed at a 10 ms hop. An acoustic path estimates how synthetic the current speech sounds using residual statistics after a light harmonic model, optionally adapted with class-conditional prototypes [4]. A linguistic path scores transcript fragments with a high-precision keyword layer and a stage machine over a vishing script. A fusion layer produces a single risk score, a coarse regime label (safe, social engineering, spoof probe, dual threat), a heuristic uncertainty band, and input-space counterfactuals obtained by re-running the same scoring function.

The contribution is not a new deep encoder. It is an evaluated, reproducible protocol for a telephony-oriented dual-stream detector, with negative results reported alongside positive ones. In particular, we show that a stage tracker helps on paraphrased scam scripts that a keyword list misses, that lightweight residual features detect a robotic pulse-formant vocoder and fail on LPC resynthesis after bandlimiting, and that disagreement-aware fusion trades false alarms for recall relative to a weighted sum.

2 Related work

The ASVspoof series is the standard public protocol for logical-access synthetic speech and physical-access replay [3, 5]. Competitive systems use self-supervised front ends and graph-attention or raw-waveform back-ends [6, 7]. Those systems are the right comparison for neural TTS detection; we do not replace them. Channel mismatch is a documented failure mode [2].

Vishing detectors in the applied literature often classify transcripts, or average an audio score with a text score. Dialogue-act tagging is a long-standing NLP task [8]; what is uncommon in telephony products is treating scam-script *progression* as a first-class streaming feature and fusing it with an acoustic authenticity score under an explicit telephone channel model.

Whole-utterance anti-spoofing misses calls that are only partially fake. Zhang et al. built PartialSpoof for that setting [9, 10]. Those systems ask whether any segment is synthetic. They do not ask whether the fake *onset* is timed to a vishing stage.

Conformal prediction gives finite-sample coverage un-

der exchangeability [11]. For streams that shift, Gibbs and Candès proposed adaptive conformal inference (ACI), a one-parameter online update of the target quantile [12]. We implement that recursion on labeled risk scores. A leftover exponential-moving-average band remains in the code as a heuristic; it is not ACI.

3 Method

3.1 Channel simulation

Incoming waveforms are resampled to 8 kHz. An optional telephone-channel simulator applies a 300–3400 Hz Butterworth bandlimit, μ -law compression with 8-bit quantization, additive noise, and packet drops with simple concealment. This is a classical DSP chain, not a learned generative model of the public switched telephone network. We use it as a stress test and, optionally, as online data augmentation.

3.2 Acoustic path

Each 25 ms speech frame is decomposed with a least-squares harmonic model. The residual is summarized by energy ratio, spectral flatness, excess kurtosis, Hilbert-envelope modulation, harmonic-to-noise ratio, phase-derivative irregularity, and an upper-band peak-to-median ratio intended to capture periodic grid energy. A 64-dimensional vector concatenates these statistics with coarse spectral descriptors.

Prototype memory maintains running means of bona fide and synthetic embeddings and scores a frame by a softmax over RMS Mahalanobis distances. The memory can be fit on labeled clips and updated with a few confirmed examples. Bootstrap Gaussians exist only so the scorer is callable before any speech is seen; they are not used in the numbers below.

3.3 Linguistic path

Transcript fragments—from an external automatic speech recognizer, or from a scheduled text source in this paper—are scored in two ways. The keyword layer is a weighted union of high-precision English patterns (gift cards, Social Security numbers, warrants). The stage tracker is a small hidden Markov model whose emissions are a broader lexicon of vishing stages (greeting, authority, problem, urgency, harvest, payment, secrecy, threat) and whose transitions prefer forward progression. The path score mixes stage depth with mean per-step log-likelihood. Isolated sensitive words in otherwise benign calls are expected to fire the keyword layer more than the stage tracker.

There is no production ASR in the experiments. The linguistic numbers therefore isolate language modeling, not recognition error.

3.4 Fusion

Let s be the acoustic synthetic probability and f the linguistic fraud probability. A naive baseline is $0.4s + 0.6f$. ShieldCall instead uses a confidence-weighted blend, a small bonus when both streams have recently been high, a social-engineering floor when f is high and s is not, and a spoof-probe floor when s is high and f is not. Counterfactual explanations zero f or s (or flatten escalation) and re-evaluate the same function. This is rule-based fusion, not causal identification.

3.5 Stage-aligned production change

Page’s CUSUM [13] watches s_t online for a mean shift. Offline we also take the two-window mean-shift time $\hat{\tau}$ that maximises $|\text{mean}_{\text{after}} - \text{mean}_{\text{before}}|$. Harvest, payment, threat, and secrecy stage times $\{t_k\}$ come from the discourse tracker. Coupling is

$$C = \max_k \exp(-(t_k - \hat{\tau})^2 / (2\sigma^2)), \quad \sigma = 0.45 \text{ s.}$$

A Monte Carlo p -value uses circular shifts of $\hat{\tau}$ on the call duration (the shift-null, not a causal graph). The operational score is $C(1 - p)$. The hypothesis is that a human-to-vocoder *handoff* aligned with a harvest stage yields larger C than an equal-duration vocoded window placed far from that stage. If mean s is matched across those two conditions, utterance-level scores cannot separate them.

3.6 Adaptive conformal inference

On a labeled stream, residuals $|s - y|$ feed ACI [12]: $\alpha_{t+1} = \Pi_{[\alpha_{\min}, \alpha_{\max}]}(\alpha_t + \gamma(\alpha - \text{err}_t))$, with interval half-width the empirical $(1 - \alpha_t)$ -quantile of past residuals. We report empirical coverage, not a finite-sample theorem for dependent telephony.

4 Experiments

All numbers in this section come from `python scripts/run_paper_experiments.py` in the accompanying repository. Sine-wave proxies are used only in unit tests and are not reported as scientific results. Mini LibriSpeech `dev-clean-2` [14] was downloaded from OpenSLR. ASVspoof was not present; the loader is implemented but unused.

4.1 Linguistic scripts

We wrote 86 English call scripts (48 train, 38 held-out) spanning impersonation of tax, benefits, banks, helpdesks, relatives, and utilities, plus benign reminders and *traps*: benign calls that mention a bank, a password reset, or a gift-card promotion without a vishing trajectory. Held-out scam lines are paraphrases that avoid many of the high-precision keywords.

Table 1: Linguistic detection. Isolated-keyword traps report mean fraud probability (lower is better).

Condition	n	EER	AUC
Train, keywords	48	0.104	0.929
Train, keywords+stages	48	0.000	1.000
Held-out, keywords	38	0.583	0.417
Held-out, keywords+stages	38	0.106	0.883
Held-out traps (mean score)	8	—	0.234




Figure 1: Linguistic AUC. The stage tracker is redundant on canonical training lines and necessary on held-out paraphrases.

Table 1 and Fig. 1 show the intended pattern. On canonical training language the keyword list is already strong. On held-out paraphrases it falls below chance (AUC 0.42), while adding the stage tracker recovers AUC 0.88. Mean score on held-out traps is 0.23. These scripts are not a sample of real victims; they test whether the stage machine does anything the keyword list does not.

4.2 Acoustic conditions

Bona fide audio is speaker-disjoint Mini LibriSpeech downsampled to 8 kHz (5 train speakers, 5 test speakers, two utterances each). Spoof audio is the same utterance passed through a vocoder. Prototype memory is fit on the training speakers only.

A pulse-formant vocoder (impulse train at estimated F_0 through three analog-style formants) is an easy condition: EER 0.00 and AUC 1.00 both clean and after 300–3400 Hz bandlimiting, on $n=20$ clips. That result shows the residual front end can fire on robotic excitation after telephone filtering. It does *not* show neural-TTS detection.

A low-order LPC vocoder on the same protocol is a hard condition. In a prior run of the same code, clean LPC yielded EER 0.40 / AUC 0.64, and narrowband LPC was at chance (EER 0.45 / AUC 0.49). Few-shot adaptation of a pulse-formant memory toward LPC reduced unseen LPC EER only from 0.50 to 0.45 after five shots. Lightweight residual features, as implemented here, are

Table 2: Operational fusion, $n=40$. Disagreement recall and safe-cell FPR use a fixed threshold of 0.5.

System	EER	AUC	Disc. rec.	Safe FPR
Acoustic only	0.267	0.833	0.90	0.80
Linguistic only	0.233	0.827	0.05	0.00
Naive sum	0.100	0.973	0.30	0.00
Logistic	0.083	0.993	1.00	0.50
ShieldCall fusion	0.000	1.000	1.00	0.40




Figure 2: Fusion on operational threat labels. ShieldCall and logistic regression recover disagreement cells that a naive sum misses, and both pay in false positives on safe calls at the 0.5 threshold.

not a substitute for ASVspoof-trained neural countermeasures [6].

4.3 Operational fusion

A call is labeled positive if it is a scam script *or* pulse-formant vocoded, and negative only if it is bona fide speech with a benign script. That is the operational question (should an operator look?) rather than the acoustic question (is the voice synthetic?). We compare acoustic-only, linguistic-only, naive weighted sum, logistic regression on (s, f) , and ShieldCall fusion. Logistic regression is fit on a grid of synthetic probabilities crossed with training-script fraud scores, not on test audio.

Table 2 and Fig. 2 are the fusion result. Linguistic-only almost never catches a vocoded call that uses mild language (disagreement recall 0.05). Naive sum improves overall AUC but still misses most disagreement cells at 0.5 (recall 0.30). ShieldCall fusion and logistic regression both reach disagreement recall 1.00. ShieldCall has perfect ranking AUC on this small set and a safe-cell false-positive rate of 0.40 at the same threshold, versus 0.00 for the naive sum. The disagreement rules buy recall; they are not free.

4.4 Timing: synthetic processes versus vocoded splices

On synthetic point processes (stage times at 5 and 6.2s, aligned alarms near 5s, unaligned near 1.1s, $n=40$), SAPC pair-win rate and AUC are both 1.00. That only shows the statistic is ordered when the inputs are clean.

On Mini LibriSpeech we splice a pulse-formant window of relative length 0.34 either at the harvest fraction or at the opposite end of a ~ 10 s stitched call ($n=6$ speaker pairs, matched mean s 0.58 vs 0.57). Coupling AUC for aligned vs unaligned is 0.47; mean s AUC is 0.53; pair-win rate is 0.50. The mix match holds; the timing claim does not. We do not tune further to force a pass.

4.5 ACI coverage

On a 400-step synthetic stream whose positive prevalence jumps from 0.25 to 0.75 at the midpoint, ACI with target $1 - \alpha = 0.90$ has empirical coverage 0.885. A frozen 90% quantile fit on the first forty residuals has coverage 0.931 on the remainder. ACI tracks the target; it is not shown to beat a frozen quantile on this particular shift.

5 Limitations

The linguistic corpus is author-written English. It does not estimate prevalence, ASR robustness, or non-English vishing. The pulse-formant condition is easy by construction. LPC and, by extension, modern neural vocoders under telephony remain unsolved for this front end. Sample sizes are small ($n \leq 48$ linguistic, $n=20$ acoustic, $n=40$ fusion). There is no user study of the explanations, no carrier deployment, and no ASVspoofer number in this paper. ACI coverage is empirical on a synthetic labeled stream. SAPC does not detect vocoded handoffs on the LibriSpeech construction we ran. Anyone citing this work as “SOTA deepfake detection” or as a demonstrated mid-call handoff detector would be misquoting it.

6 Conclusion

ShieldCall is a streaming dual-stream detector for telephone-bandwidth audio and call transcripts. The result that is worth reproducing is not a leaderboard EER. It is that a scam-stage tracker still works when a keyword list is paraphrased away, that residual acoustics can detect a robotic vocoder after bandlimiting and fail on LPC, and that disagreement-aware fusion changes the recall/false-alarm tradeoff in the direction an operator would expect. The next measurements that would change the claim are ASVspoofer Logical Access through the same channel simulator, and a handoff test on PartialSpoofer-style splices with real transcripts where SAPC either succeeds or fails again in public.

Reproducibility

Code, configs, scripts, and the Mini LibriSpeech downloader are in the companion repository. The experiment command is `python scripts/run_paper_experiments.py`. Set `SHIELD_CALL_ASVSPOOF_ROOT` to enable the unused ASVspoofer loader.

References

- [1] Federal Trade Commission. Consumer sentinel network data book. Technical report, FTC, 2024.
- [2] You Zhang, Fei Jiang, and Zhiyao Duan. An empirical study on channel effects for synthetic voice spoofing countermeasure systems. *arXiv preprint arXiv:2104.01320*, 2021.
- [3] Massimiliano Todisco, Xin Wang, Ville Vestman, Md Sahidullah, Héctor Delgado, Andreas Nautsch, Junichi Yamagishi, Nicholas Evans, Tomi Kinnunen, and Kong Aik Lee. ASVspoofer 2019: Future horizons in spoofed and fake audio detection. In *Proc. Interspeech*, pages 1008–1012, 2019.
- [4] Jake Snell, Kevin Swersky, and Richard Zemel. Prototypical networks for few-shot learning. In *NeurIPS*, 2017.
- [5] Xin Wang, Junichi Yamagishi, Massimiliano Todisco, Héctor Delgado, Andreas Nautsch, Nicholas Evans, Md Sahidullah, Ville Vestman, Tomi Kinnunen, Kong Aik Lee, et al. ASVspoofer 2019: A large-scale public database of synthesized, converted and replayed speech. *Computer Speech & Language*, 64:101114, 2020.
- [6] Jee-weon Jung, Hee-Soo Heo, Hemlata Tak, Hye-jin Shim, Joon Son Chung, Bong-Jin Lee, Ha-Jin Yu, and Nicholas Evans. AASIST: Audio anti-spoofing using integrated spectro-temporal graph attention networks. In *ICASSP*, 2022.
- [7] Hemlata Tak, Jose Patino, Massimiliano Todisco, Andreas Nautsch, Nicholas Evans, and Anthony Larcher. End-to-end anti-spoofing with RawNet2. In *ICASSP*, 2021.
- [8] Andreas Stolcke, Klaus Ries, Noah Coccaro, Elizabeth Shriberg, Rebecca Bates, Daniel Jurafsky, Paul Taylor, Rachel Martin, Carol Van Ess-Dykema, and Marie Meteer. Dialogue act modeling for automatic tagging and recognition of conversational speech. *Computational Linguistics*, 26(3):339–373, 2000.
- [9] Lin Zhang, Xin Wang, Erica Cooper, Junichi Yamagishi, Jose Patino, and Nicholas Evans. An initial investigation for detecting partially spoofed audio. In *Proc. Interspeech*, 2021.

- [10] Lin Zhang, Xin Wang, Erica Cooper, Nicholas Evans, and Junichi Yamagishi. The PartialSpoof database and countermeasures for the detection of short fake speech segments embedded in an utterance. *IEEE/ACM Trans. Audio, Speech, Lang. Process.*, 31:813–825, 2023.
- [11] Anastasios N. Angelopoulos and Stephen Bates. A gentle introduction to conformal prediction and distribution-free uncertainty quantification. *Foundations and Trends in Machine Learning*, 2023.
- [12] Isaac Gibbs and Emmanuel Candès. Adaptive conformal inference under distribution shift. In *Advances in Neural Information Processing Systems*, 2021.
- [13] E. S. Page. Continuous inspection schemes. *Biometrika*, 41(1/2):100–115, 1954.
- [14] Vassil Panayotov, Guoguo Chen, Daniel Povey, and Sanjeev Khudanpur. LibriSpeech: An ASR corpus based on public domain audio books. In *ICASSP*, pages 5206–5210, 2015.